

PAGANI: A PARALLEL ADAPTIVE GPU ALGORITHM FOR NUMERICAL INTEGRATION



Authors: Ioannis Sakiotis, Kamesh Arumugam, Marc Paterno, Desh Ranjan, Balsa Terzic, Mohammad Zubair

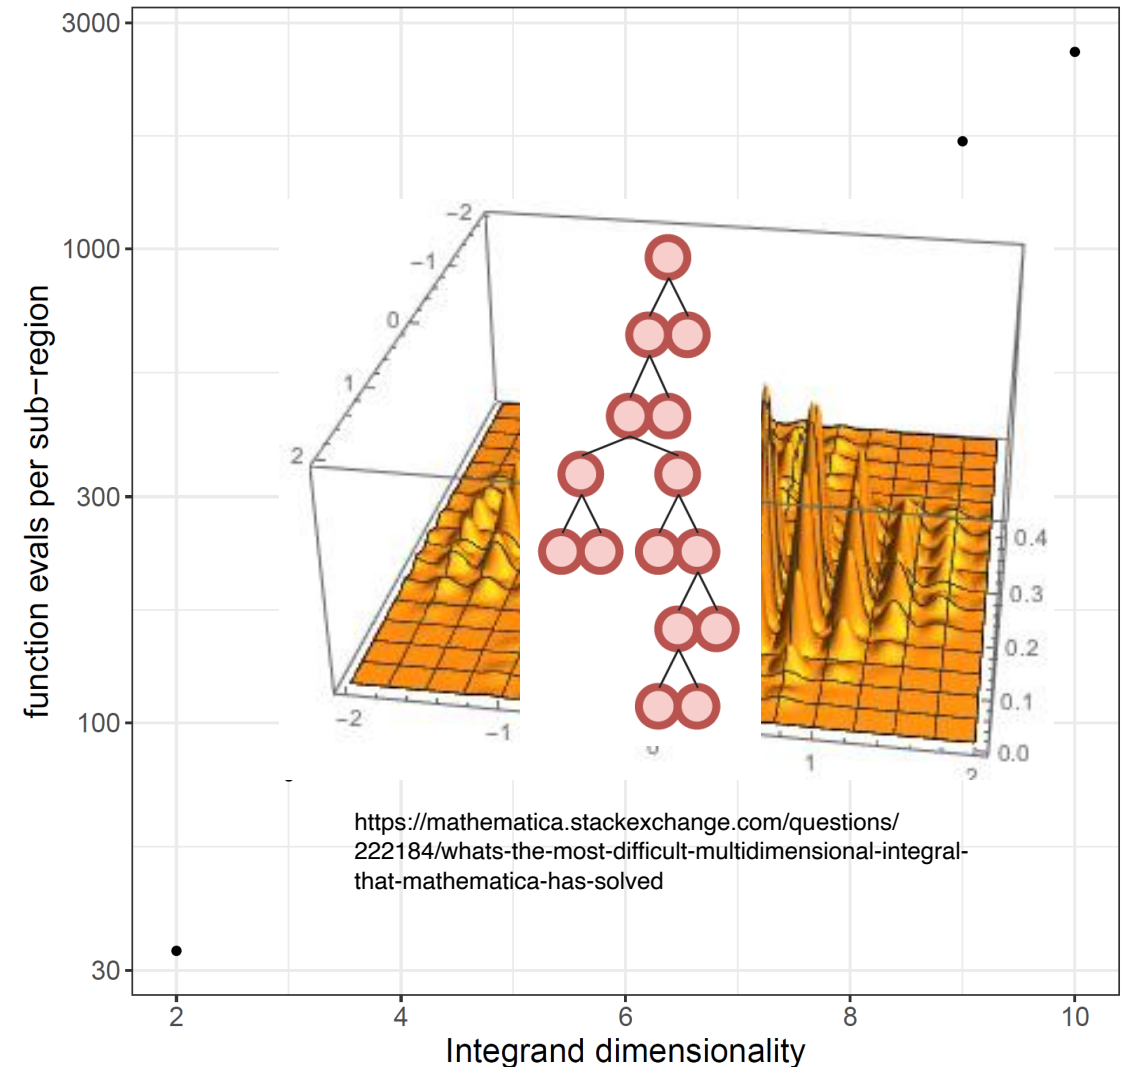


INTRODUCTION

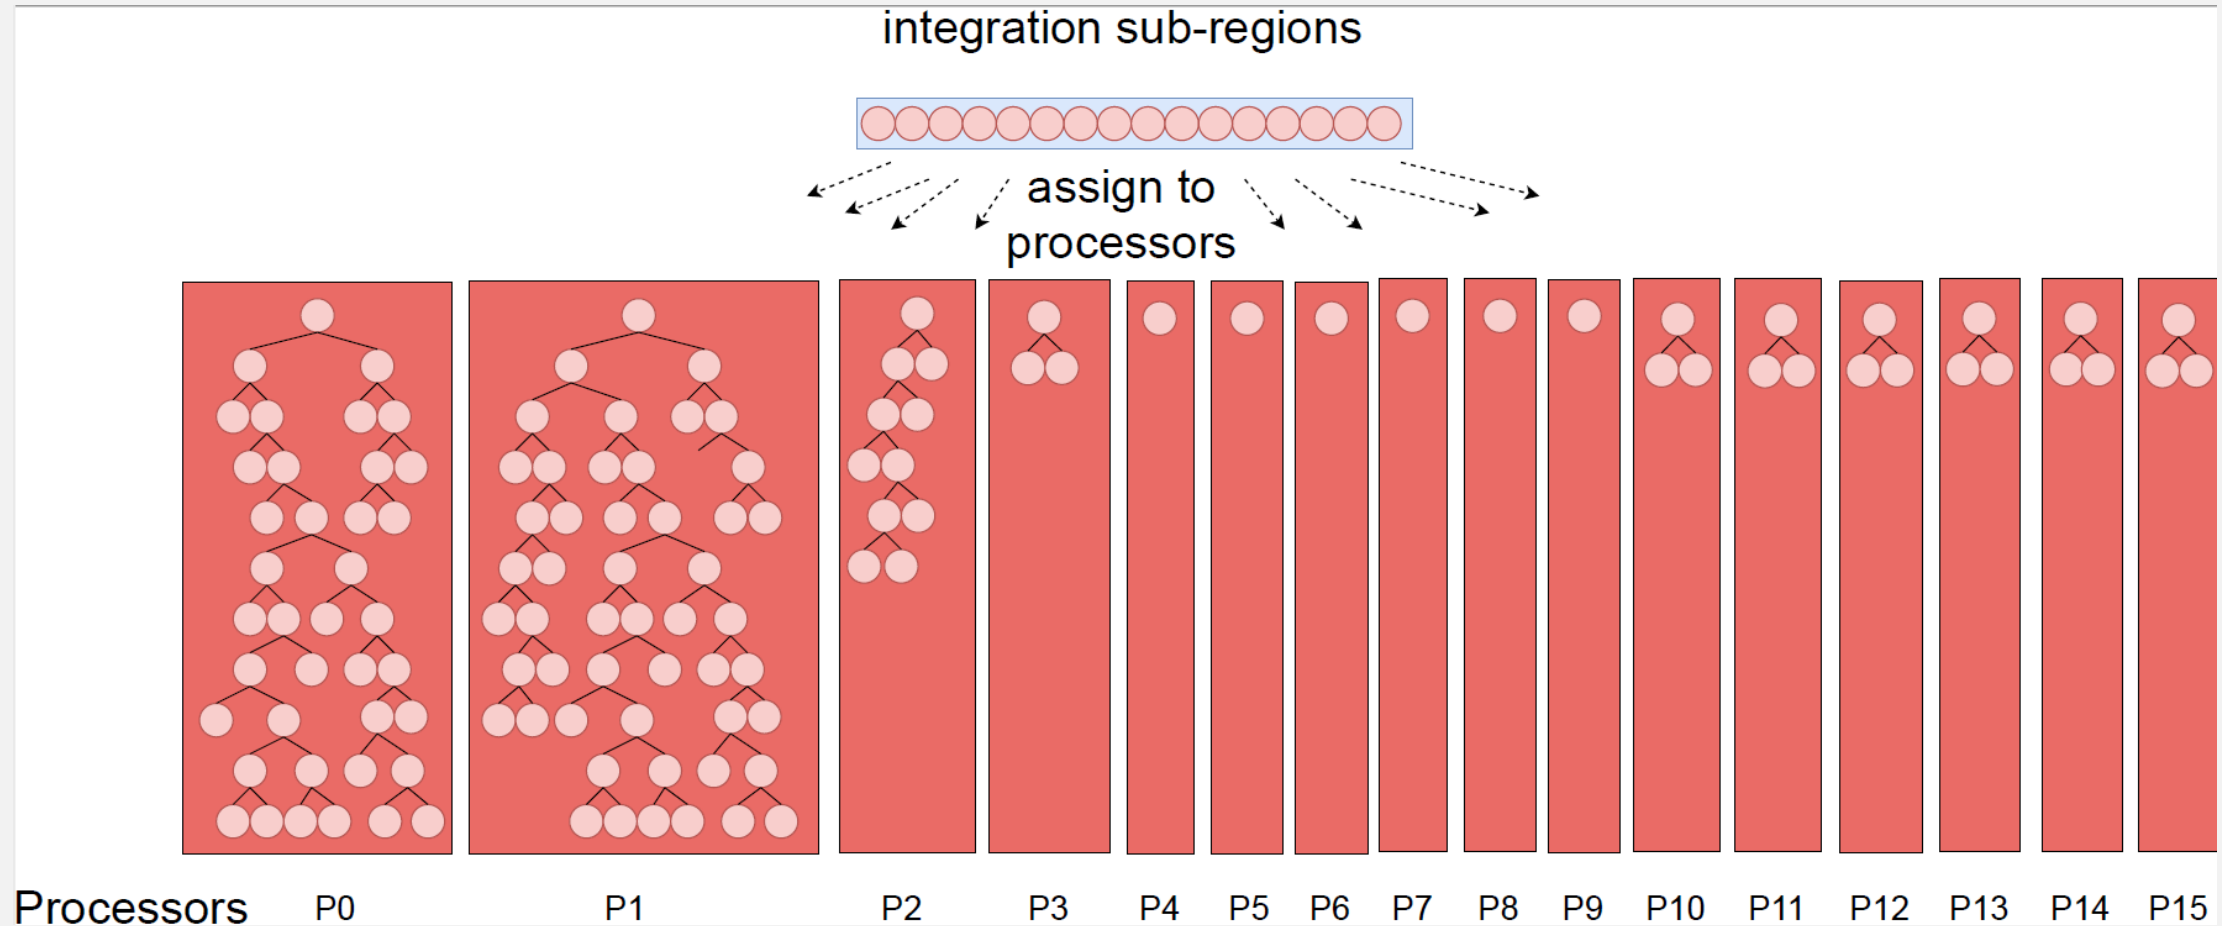
- Applications: parameter estimation, simulation of beam dynamics, risk management, ray-tracing
- Probabilistic, deterministic
- No algorithm guarantees the accuracy of its estimated results
- Evaluate accuracy
- Specify desired accuracy (relative error tolerance, digits of precision, etc.)
- Goal: Bring adaptive quadrature to GPUs

BACKGROUND: QUADRATURE

- Weighted summation $\sum (w_i f(x_i))$
- Error-estimate
- Expect large initial error
- Apply weighted summation in sub-region to reduce error
- Highly parallelizable
- Number of points grows exponentially with the number of dimensions
- Uniform split and sub-region evaluation infeasible
- Cuhre
 - $2^n + \Theta(n^3)$ functions evaluations (n-dimensions)
 - Attractive option for low/mid dimensional integrands
 - Priority-queue




NAÏVE PARALLELIZATION



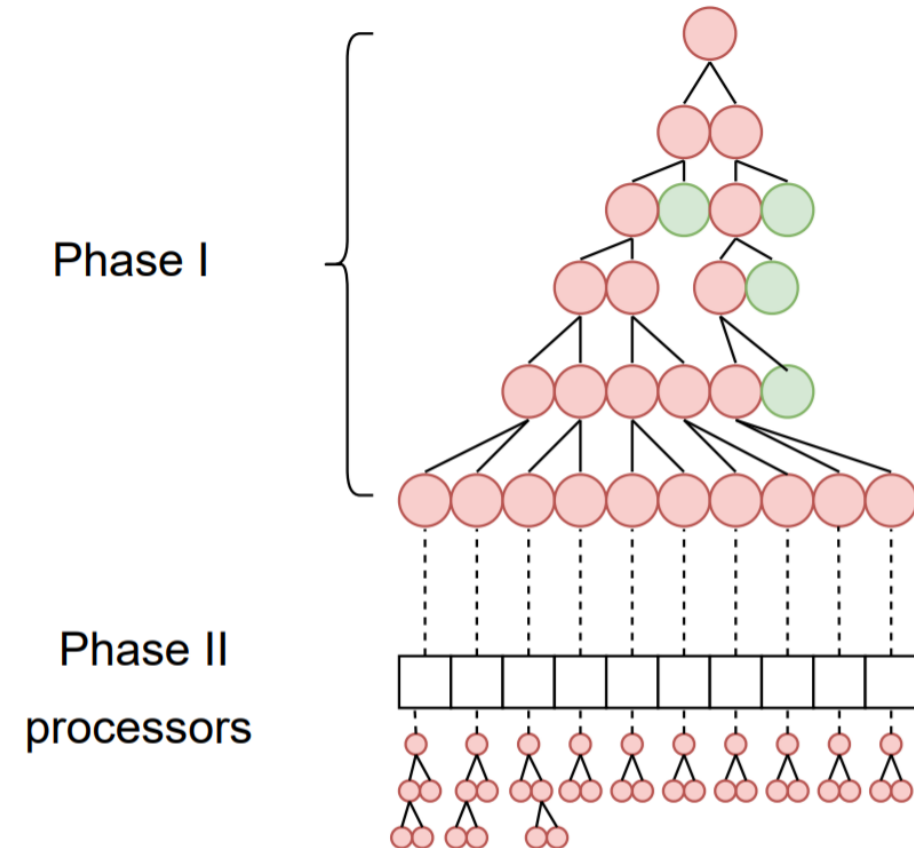
BACKGROUND: TWO-PHASE CUHRE

- GPU-targeted algorithm
- Assign each sub-region to a processor
- **Applies the Cuhre algorithm on each sub-region in parallel**
- Utilizes pre-processing Phase 1
 - Generate sufficiently large workload (# of sub-regions)
 - Load-balancing
- No synchronization between processors
- Local termination



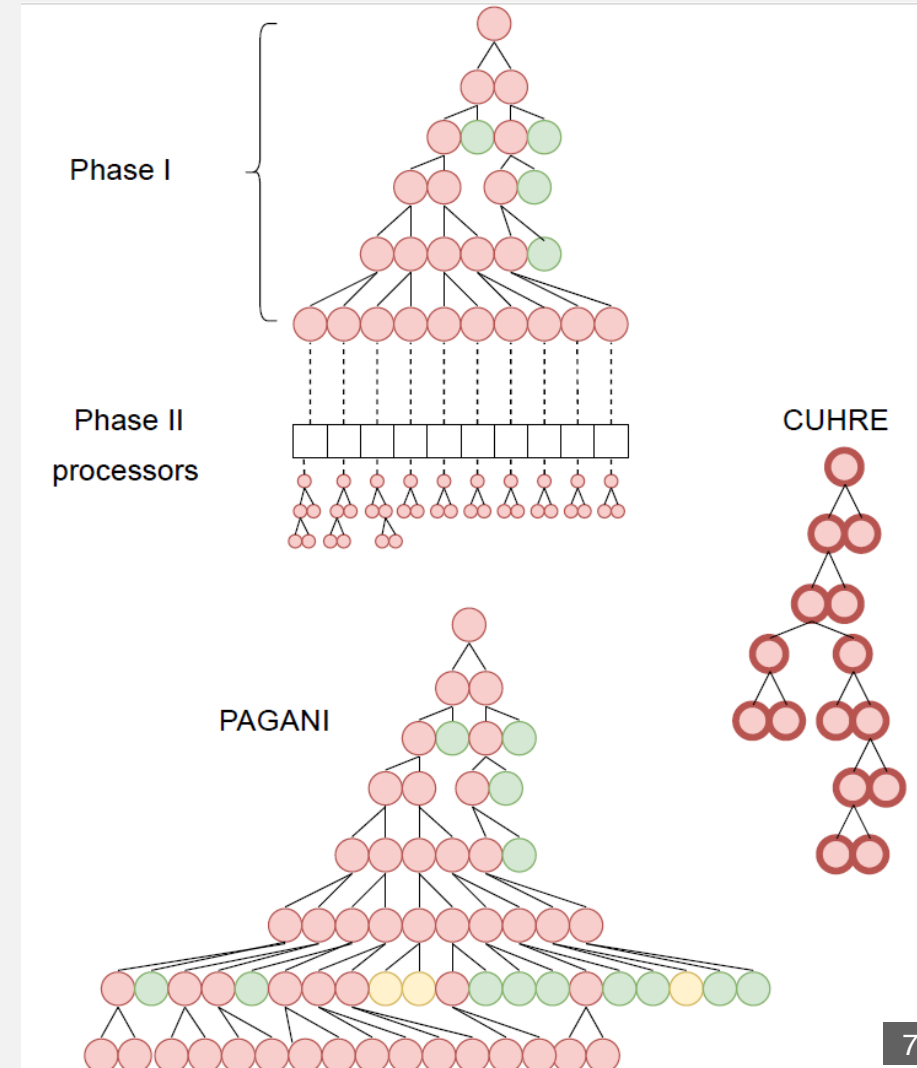
TWO-PHASE CUHRE

- Utilization of sequential algorithm by parallel processors
- Poor load-balancing when high-precision
- Local data-structures
- Unknown global state (unless global sync.)
- Local termination



PAGANI

- Parallel algorithm designed for massively parallel architectures
- Avoid sequential scheme
- Sub-divide all sub-regions
- Filtering instead of sorting
- Green/yellow = accurate enough
- Avoid synchronization after isolated processing
- Uniform workload
- Bound by memory



ALGORITHM DESCRIPTION

- Initial uniform-split
- Parallel sub-region evaluation
- Two-level error-estimate
- Relative error classification
 - Finished/active
- Summations
- Termination conditions
- Conditional threshold classification
- Filtering
- Split all active regions

Algorithm 2 PAGANI Algorithm

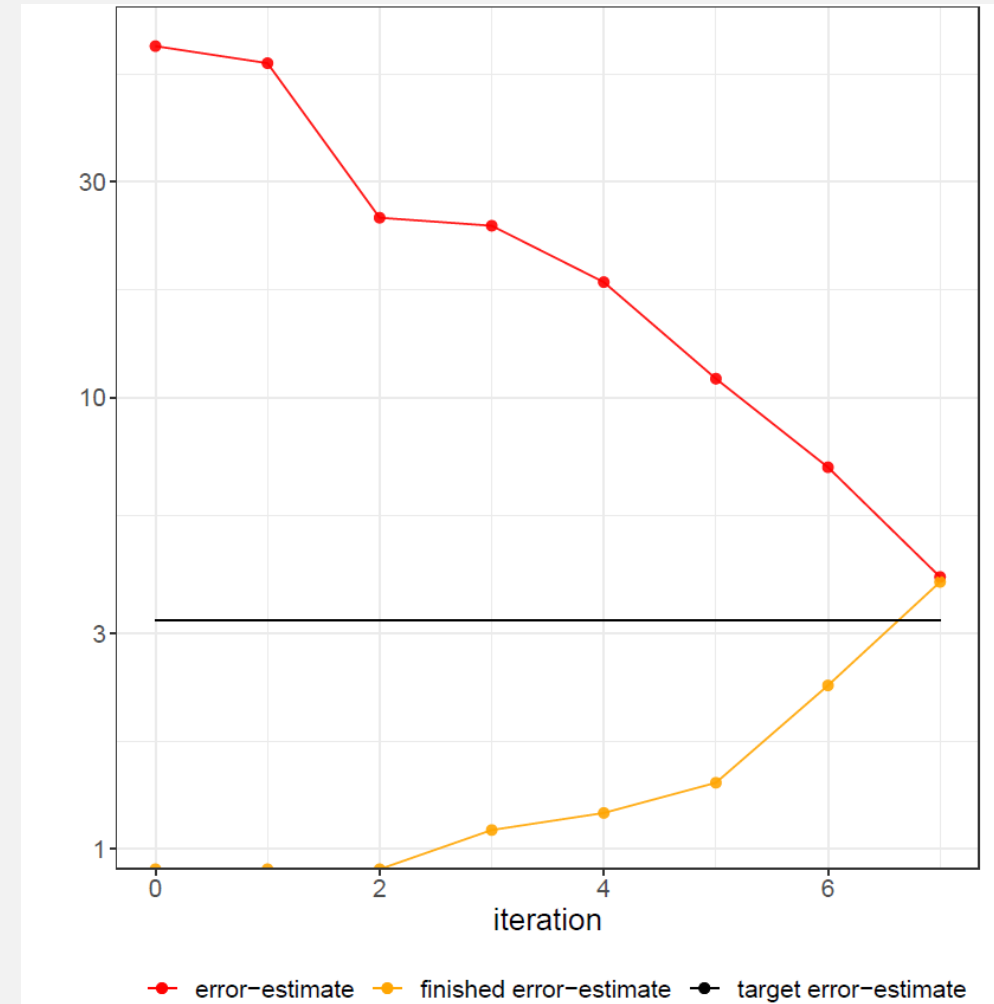
```

1: procedure PAGANI( $f, n, b[n], \tau_{rel}, \tau_{abs}$ )
2:    $R_0 \leftarrow b$ 
3:    $s \leftarrow d^n$  ▷ region list size
4:    $H \leftarrow \text{UNIFORM-SPLIT}(R_0, d)$ 
5:    $A[1 : s] \leftarrow 1$ 
6:    $V[1 : s], E[1 : L], K[1 : s] \leftarrow 0$ 
7:    $V_p[1 : s], E_p[1 : s] \leftarrow 0$ 
8:    $v, e, v_f, e_f \leftarrow 0$  ▷ cumulative/finished estimates
9:   for  $it : it_{max}$  do
10:     $V, E, K \leftarrow \text{EVALUATE}(H)$ 
11:     $E \leftarrow \text{TWO-LEVEL-ERROR}(V, E, V_p, E_p)$ 
12:     $A \leftarrow \text{REL-ERR-CLASSIFY}(V, E, A)$ 
13:     $v \leftarrow \text{SUM}(V)$ 
14:     $e \leftarrow \text{SUM}(E)$ 
15:    if  $\frac{e+e_f}{|v+v_f|} \leq \tau_{rel}$  or  $e + e_f \leq \tau_{abs}$  then
16:      return  $v + v_f, e + e_f$ 
17:     $A \leftarrow \text{THRESHOLD-CLASSIFY}(A, E, v + v_f, e + e_f, v, e, s)$ 
18:     $v_f \leftarrow v - \text{SUM}(V \cdot A) + v_f$ 
19:     $e_f \leftarrow e - \text{SUM}(E \cdot A) + e_f$ 
20:     $H, V, E, L \leftarrow \text{FILTER}(H, V, E, A)$ 
21:     $V_p \leftarrow V, E_p \leftarrow E$  ▷ update all parents
22:     $H \leftarrow \text{SPLIT}(H, K)$ 
23:     $s \leftarrow 2s$ 

```


SUB-REGION CLASSIFICATION

- Why classify finished/active?
 - Only keep active regions in memory
 - Use relative error for global termination and finished/active classification
 - Regions with small estimates may not satisfy relative error termination
 - Remove regions that don't contribute "significantly"
- Aggressive filtering
 - How to define "significantly"?
 - Finished regions irrecoverable for performance
 - Balance finished and active estimates
 - Pick a threshold (initially the average) and adapt until criteria are met
 - Criteria: conserved memory, finished vs. active ratio
 - Perform if memory exhaustion or convergence of significant digits



INTEGRAND TEST SUITE

- First six integrands represent challenging integrand families
- Product and corner peaks, oscillatory, Gaussian, etc.
- Typically, randomized parameters
- Fixed parameters

$$f_1(x) = \cos\left(\sum_{i=1}^8 i x_i\right)$$

$$f_2(x) = \prod_{i=1}^6 \left(\frac{1}{50^2} + (x_i - 1/2)^2\right)^{-1}$$

$$f_3(x) = \left(1 + \sum_{i=1}^d i x_i\right)^{-d-1}$$

$$f_4(x) = \exp\left(-625 \sum_{i=1}^d (x_i - 1/2)^2\right)$$

$$f_5(x) = \exp\left(-10 \sum_{i=1}^d |x_i - 1/2|\right)$$

$$f_6(x) = \begin{cases} \exp\left(\sum_{i=1}^6 (i+4) x_i\right) & \text{if } x_i < (3+i)/10 \\ 0 & \text{otherwise} \end{cases}$$

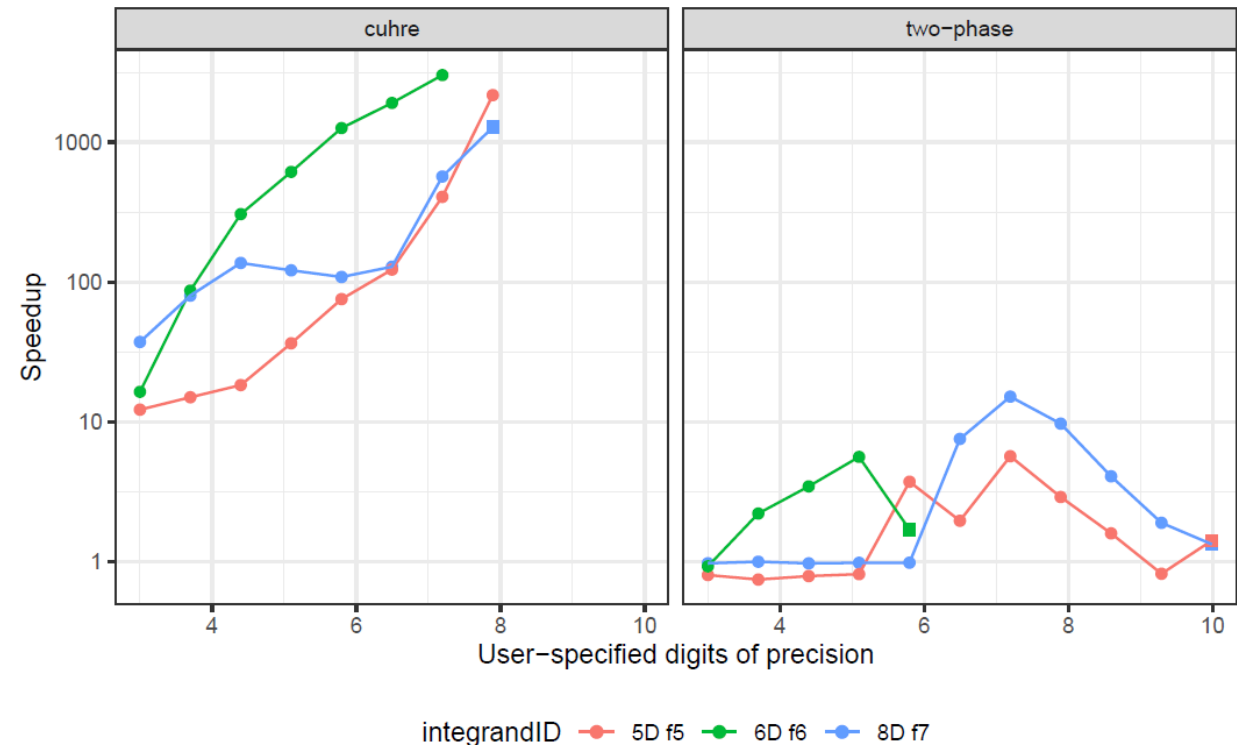
$$f_7(x) = \left(\sum_{i=1}^d x_i^2\right)^{11}$$

$$f_8(x) = \left(\sum_{i=1}^d x_i^2\right)^{15/2}$$

EXPERIMENTAL RESULTS

- CUDA 11 implementation
- Executed on V100 16 GB
- 2.4 GHz Xeon R Gold 6130 CPU
- Orders of magnitude speedup over sequential Cuhre
- Improved robustness over two-phase
 - Improved load-balancing on higher precision
 - More reliable error-estimate
- Comparable performance on low-precision

3 digits of precision = $1e-3$ relative-error tolerance
4 digits of precision = $1e-4$ relative-error tolerance

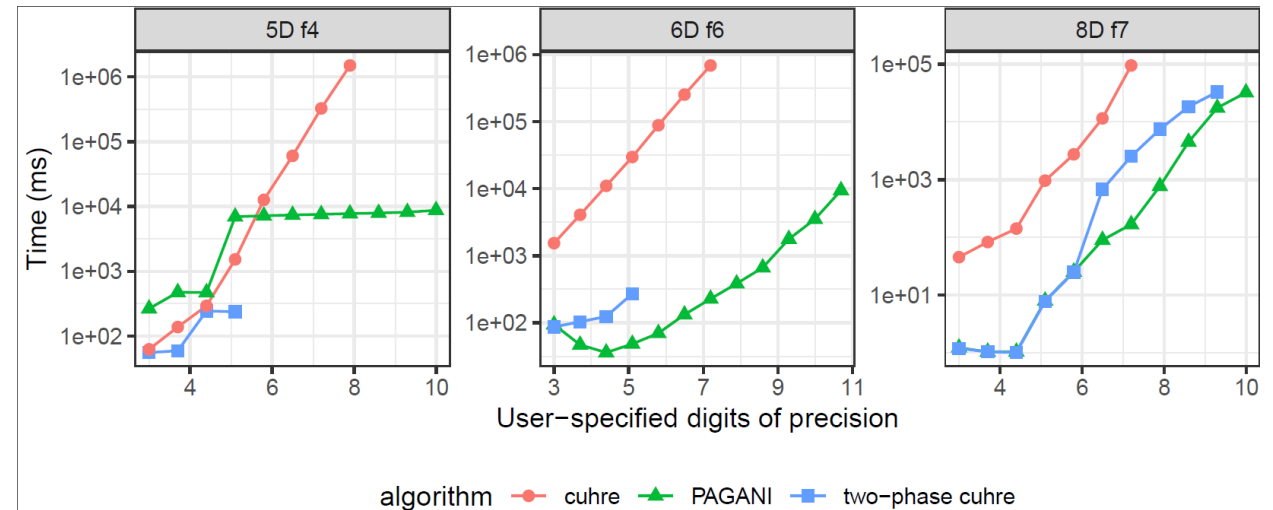


PERFORMANCE

Pagani generates more sub-regions

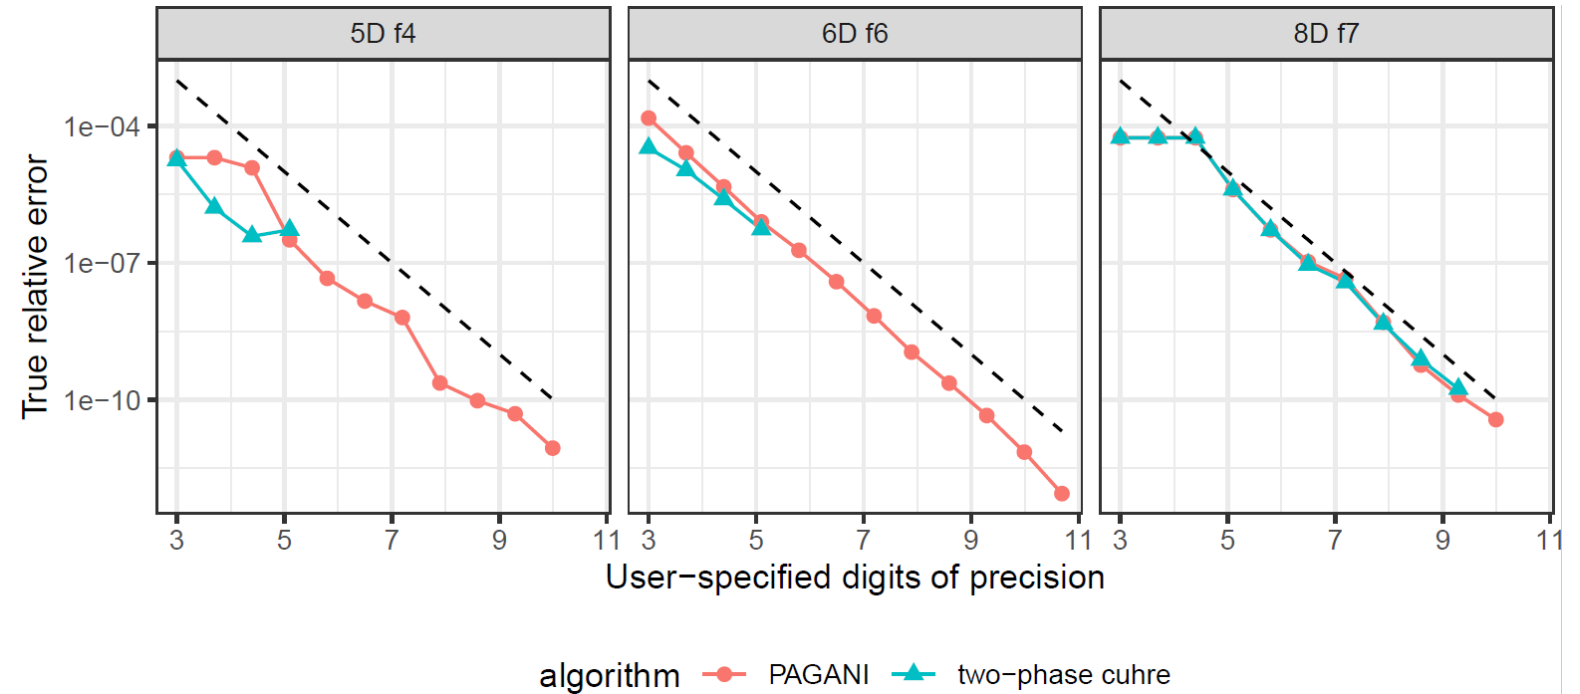
“Easy” integrals not worth the overhead

More digits-of-precision




Accuracy

- No randomized parameters
- Evaluate error-estimate
- Claimed vs true relative-error
- Proximity to line indicates sub-division efficiency
- Check bounds of error estimate




ACCURACY

- Estimates as measurement
- Consistency across iterations
- Overlapping error-bars
- 3-digits of precision
- Red line = true value



ACCURACY

- Estimates as measurement
- Consistency across iterations
- Overlapping error-bars
- 6-digits of precision
- Red line = true value
- More iterations



CONCLUSION

- New deterministic adaptive algorithm for highly parallel architectures
- No use of sequential algorithm
- Orders-of-magnitude speedup on challenging integrals
- Comparable performance with Two-Phase on low-precision
- 4-15 speedup over Two-Phase on medium/high precision
- Improved robustness and execution time
- Reliable error-estimation
- At least as reliable as Cuhre

FUTURE WORK

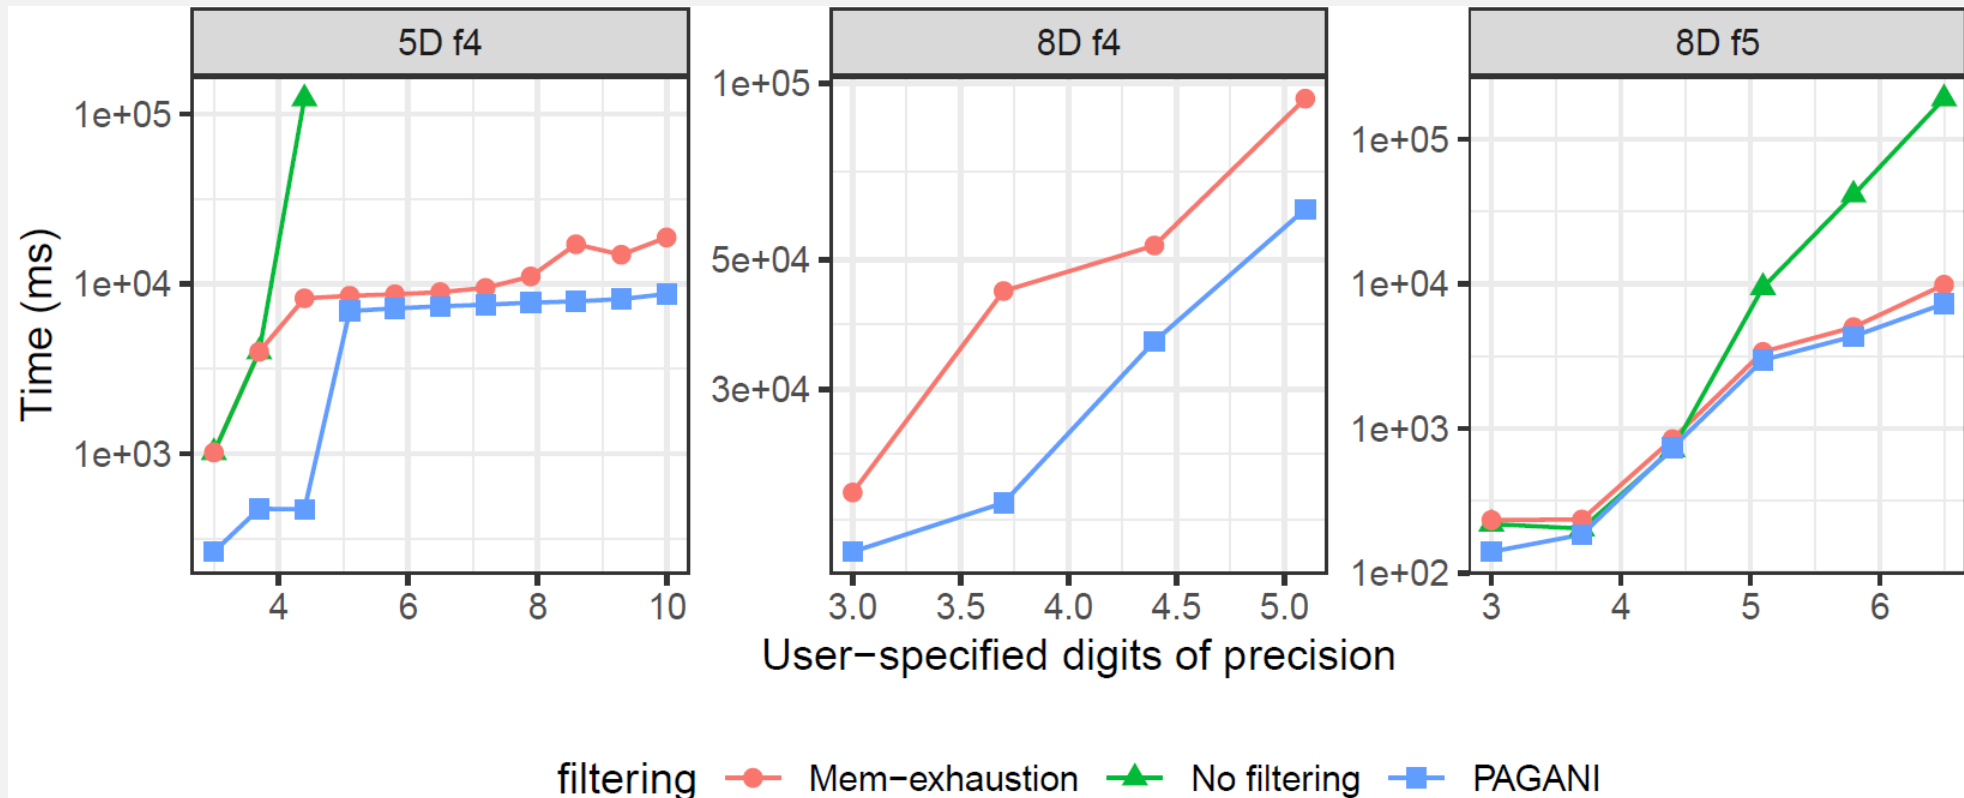
- Ampere architecture
- Multi-GPU
- Kokkos: 15% overhead in main computational kernels

QUESTIONS

IMPORTANT ALGORITHM CHARACTERISTICS

- All operations utilize parallelization
 - Function-evaluations within sub-region
 - Sub-region evaluations
- No use of sequential algorithm
- No persistence in region-processor mapping
- Global data-structure
- Global state through reduction
- Implicit synchronization

PERFORMANCE: HEURISTIC SEARCH FILTERING




APPENDIX

BACKGROUND: CUHRE

- Maintains priority list based on highest error-estimate
- Sub-divides sub-region with highest error-estimate
- Sum integral estimate of all sub-regions to yield final integral estimate
- “Harder” integrals require more sub-regions
- Disjoint sub-region are independent

Algorithm 1 Sequential Adaptive Integration Algorithm

```
1: procedure ADAPTIVE NUMERICAL INTEGRATION( $f, b[n]$ )
2:    $H \leftarrow b$ 
3:   while (termination condition is not satisfied) do
4:     Extract a non-empty subset of regions  $S$  from  $H$ 
5:     for each region  $R \in S$  do
6:       partition  $R$  into  $k$  regions along split-axis
7:       Let  $R_1, R_2 \dots R_k$  be these  $k$  regions
8:       for  $i \leftarrow 1 \dots k$  do
9:          $IntEst(R_i) \leftarrow$  Integral estimate for  $R_i$ 
10:         $ErrEst(R_i) \leftarrow$  Error estimate for  $R_i$ 
11:         $axis_i \leftarrow$  split-axis for  $R_i$ 
12:        Insert  $R_i$  into  $H$ 
13:       $GlobalIntEst \leftarrow \sum_{R_i \in H} IntEst(R_i)$ 
14:       $GlobalErrEst \leftarrow \sum_{R_i \in H} ErrEst(R_i)$ 
15:    return ( $GlobalIntEst, GlobalErrEst$ )
```



RELATIVE ERROR CLASSIFICATION

LEMMA 3.1. *Let m denote the number of subregions. Assume that $\forall i \ 1 \leq i \leq m$, $e_i \geq 0$ and v_i 's have the same sign. Let $e = \sum_{i=1}^m e_i$ denote the cumulative error and $v = \sum_{i=1}^m v_i$ denote the cumulative integral estimate. Suppose $\forall i \ 1 \leq i \leq m$, $e_i \leq |v_i| \cdot \tau_{rel}$. Then $e \leq |v| \cdot \tau_{rel}$.*

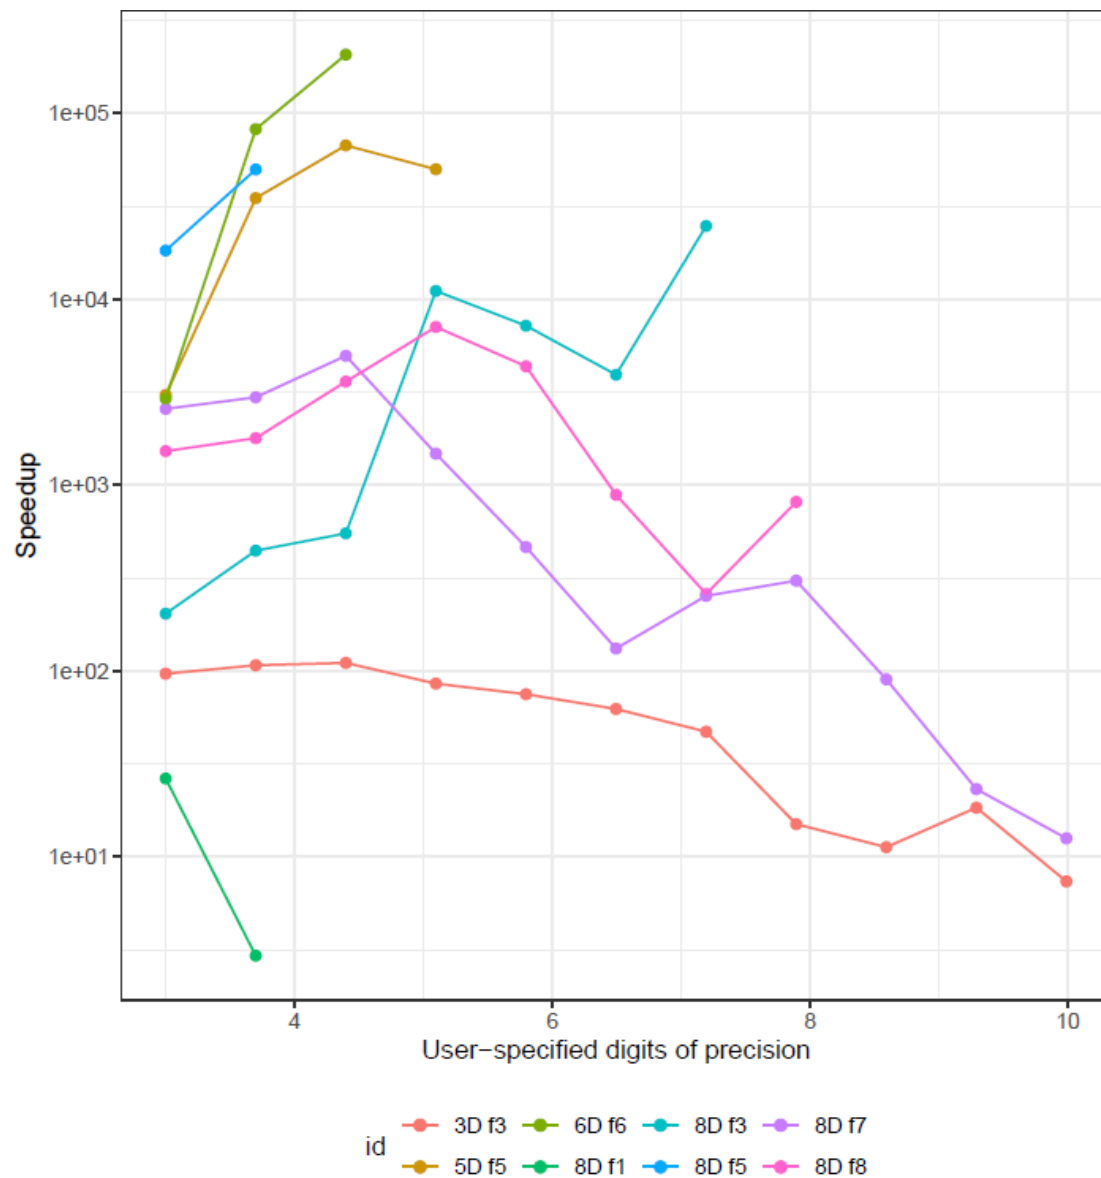
PROOF. $e = \sum_{i=1}^m e_i \leq \sum_{i=1}^m |v_i| \cdot \tau_{rel} \leq |\sum_{i=1}^m v_i| \cdot \tau_{rel} = |v| \cdot \tau_{rel}$. \square

PURE ERROR ESTIMATE

- Error coefficients
- If $coeff[0] * rule[1] \leq rule[2]$ $coeff[0] * rule[2] \leq sum[3]$
 - Return $coeff[1] * rule[1]$
 - Otherwise return $coeff[2] * \max(\max(rule[1], rule[2]), rule[3])$

TWO-LEVEL ERROR ESTIMATE

- Pure error-estimate
- Integral estimate
- Sibling estimate
- Parent estimate
- Diff: $\left| .25(\text{siblRes} + \text{selfRes} - \text{parRes}) \right|$
- $C = 1 + 2\text{diff}/(\text{selfErr} + \text{siblErr})$
- $\text{selfErr} = c * \text{selfErr}$
- $\text{selfErr} == \text{selfErr} + \text{diff}$



SPEEDUP
OVER
QUASI-
MONTE
CARLO

PERFORMANCE BREAKDOWN

- Sub-region Evaluation: more than **90%** of execution time
 - 40% peak-performance for double-precision
 - Compute-bound
 - Limited number of memory accesses (initial read, final write)
- Post-processing, classification, filtering, sub-division
 - Thrust library
- All operations on GPU

Algorithm 2 PAGANI Algorithm

```

1: procedure PAGANI( $f, n, b[n], \tau_{rel}, \tau_{abs}$ )
2:    $R_0 \leftarrow b$ 
3:    $s \leftarrow d^n$  ▷ region list size
4:    $H \leftarrow \text{UNIFORM-SPLIT}(R_0, d)$ 
5:    $A[1 : s] \leftarrow 1$ 
6:    $V[1 : s], E[1 : L], K[1 : s] \leftarrow 0$ 
7:    $V_p[1 : s], E_p[1 : s] \leftarrow 0$ 
8:    $v, e, v_f, e_f \leftarrow 0$  ▷ cumulative/finished estimates
9:   for  $it : it_{max}$  do
10:     $V, E, K \leftarrow \text{EVALUATE}(H)$ 
11:     $E < - \text{TWO-LEVEL-ERROR}(V, E, V_p, E_p)$ 
12:     $A < - \text{REL-ERR-CLASSIFY}(V, E, A)$ 
13:     $v \leftarrow \text{SUM}(V)$ 
14:     $e \leftarrow \text{SUM}(E)$ 
15:    if  $\frac{e+e_f}{|v+v_f|} \leq \tau_{rel}$  or  $e + e_f \leq \tau_{abs}$  then
16:      return  $v + v_f, e + e_f$ 
17:     $A \leftarrow \text{THRESHOLD-CLASSIFY}(A, E, v + v_f, e + e_f, v, e, s)$ 
18:     $v_f \leftarrow v - \text{SUM}(V \cdot A) + v_f$ 
19:     $e_f \leftarrow e - \text{SUM}(E \cdot A) + e_f$ 
20:     $H, V, E, L \leftarrow \text{FILTER}(H, V, E, A)$ 
21:     $V_p \leftarrow V, E_p \leftarrow E$  ▷ update all parents
22:     $H \leftarrow \text{SPLIT}(H, K)$ 
23:     $s \leftarrow 2s$ 

```

Algorithm 3 Threshold Classification Algorithm

```
1: procedure THRESHOLD-CLASSIFY( $A, E, v_{tot}, e_{tot}, e_{it}, s_{it}$ )
2:    $P_{max} \leftarrow .25$   $\triangleright$  target percentage of error budget
3:    $e_{\bar{a}} \leftarrow 0$   $\triangleright$  error-estimate contribution from inactive regions
4:    $s_{\bar{a}} \leftarrow 0$   $\triangleright$  # inactive regions
5:    $min, max \leftarrow MinMax(E)$ 
6:    $e_b \leftarrow e_{tot} - v_{tot} \cdot \tau_{rel}$   $\triangleright$  error budget
7:    $t \leftarrow \frac{e_{it}}{s_{it}}$   $\triangleright$  set initial threshold as avg. error-estimate
8:   repeat
9:      $A \leftarrow APPLY-THRESHOLD(E, t)$ 
10:     $s_{\bar{a}} \leftarrow n - SUM(A)$ 
11:    if  $s_{\bar{a}} > .5 \cdot s_{it}$  then  $\triangleright$  memory requirement
12:       $e_{\bar{a}} \leftarrow e_{it} - SUM(A \cdot E)$ 
13:      if  $e_{\bar{a}} \leq P_{max} \cdot e_b$  then  $\triangleright$  accuracy requirement
14:        return  $A$ 
15:       $UPDATE-THRESHOLD(t, s_{\bar{a}}, e_{\bar{a}}, P_{max}, min, max)$ 
16:    until  $\frac{s_{\bar{a}}}{s} > .5$  and  $e_{\bar{a}} \leq P_{max} \cdot e_b$ 
17:    return  $A$   $\triangleright$  unsuccessful filtering
```

HEURISTIC THRESHOLD SEARCH

- Classify many regions with “small” contributions
- How small is small enough?
- Try a different values and observe effect before committing to filtering
- Heuristic filtering: average error-estimate as initial threshold
- Try current threshold
- If % of finished regions is too small, move towards max
- Else move towards min
- Try again

